

00:00:00,100 --> 00:00:30,300 [speaker_0]

Let's face it. Most people are still using AI like it's 2022, asking ChatGPT to write a blog post or social captions, when the real game-changer is AI agents. Yet, very few people are actually doing them right. So in this video today, I'm gonna share with you what AI agents actually are, the Agent Foundation's Playbook, which is my framework for creating AI agents, how to spot agent-ready tasks in your business, how to build your first AI agent, and also mistakes to avoid as you get into

00:00:30,440 --> 00:01:00,530 [speaker_0]

building them for your business. By the way, if you're new here, my name is Rick Mulready. And for the past 11 and a half years, I've run a multi-seven-figure online business, helping other online business owners grow and scale their business more efficiently. Today, I run a community called The AI Playbook, where I help online businesses leverage AI so that you can streamline your business, become more efficient and profitable using AI in your business. I'll link to the community in the description below. So before we dive into anything, you need to understand that there are actually three different modalities, if you will, of AI. So first, you've

00:01:00,560 --> 00:01:30,650 [speaker_0]

got chatbots. Think ChatGPT, Claude, Gemini, et cetera. Those are reactive. So you ask a question, then they give you an answer, right? The second modality is copilots, so like Cursor for coding or Microsoft Copilot. These are assistants that live inside your existing tools, helping you work faster and more efficiently within that specific environment. But here's where it gets interesting. Agents are completely different. Tools like Relay.app, n8n, Zapier, Make,

00:01:30,719 --> 00:02:00,719 [speaker_0]

Lindy.ai, they don't necessarily wait for you to ask questions. They can operate autonomously. Think of it this way. ChatGPT is like having an intern with a PhD, mind you. You give them a task, they complete it, and then they're waiting for the next assignment, if you will. But an AI agent, think of it like, that's like having James Bond on your team. You give them a mission. You give them the resources, and then they figure out how to complete it. They make the decisions. They adapt when things go wrong, and they keep

00:02:00,780 --> 00:02:30,780 [speaker_0]

working even if you're sleeping. Every true AI agent needs four foundations, and I call this the Agent Foundation's Playbook. So the first foundation is the intelligence of the AI agent. This is the brain. This is your LLM, the AI model. This could be GPT-4. This could be Claude. This could be Gemini. This is what processes information. But intelligence alone isn't enough, and that's where the second part of the foundation comes in. This is the integrations. These are the hands and feet of your

00:02:30,879 --> 00:03:01,139 [speaker_0]

agent, if you will, the ability to actually do things in your business. These are the tools that you want the agent to use, so you can connect to Slack, for example, or Google Sheets, your CRM, your email. And that leads me to foundation number three, which is

the instructions. This is the DNA of your agent. This is the system prompt or the instructions that define its personality, its rules of engagement, if you will, its boundaries. And then foundation number four is the memory. This is what separates agents from simple automations. Memory

00:03:01,180 --> 00:03:31,340 [speaker_0]

comes in two forms: short-term memory. That's like the context window, for example, re-remembering what happened in this specific conversation or task, and then long-term memory. That's learning from past interactions and building a knowledge base over time. But here's the thing. Most people focus only on intelligence, unfortunately. They connect, you know, ChatGPT-5 or Claude-4 Sonnet to something and then wonder why it's not working like an actual agent. It's like trying to build a car with just an engine. You need

00:03:31,439 --> 00:04:01,750 [speaker_0]

all four wheels to actually go somewhere. All right, now that we have the agent foundations, how do you know which parts of your business are actually ready for agents? Because obviously, not everything should be automated in your business. So here's what I want you to do. For the next three to five business days, and look, I know this is the worst exercise in the world. It's tedious, but there is gold when you finish this exercise. I want you to audit your time. Track what you're actually doing hour by hour, and you're looking for three specific patterns at

00:04:01,800 --> 00:04:31,860 [speaker_0]

the end of the three to five days. So the first pattern you're looking for is the repetitive tasks, like things you do the same way every time, like social media posting, email responses, data entry, research, for example. If you can write down the steps, an agent can probably do it. The second pattern is clear SOPs, so standard operating procedures. If you have a documented process or even if just in your head, for example, but follows the same pattern every time you do it, that's agent territory for you right there. And then

00:04:31,920 --> 00:05:02,459 [speaker_0]

pattern number three is those draining tasks, but really necessary. You know those tasks that just suck the life out of you, but they've gotta get done. And then once you've found those patterns, before you automate anything, still run it through these three filters. And the first filter is, is it important? Don't automate busy work that shouldn't exist in the first place, right? Fix the process first, and then automate it. The second filter is, can I map the steps? If you can't explain exactly

00:05:02,500 --> 00:05:32,514 [speaker_0]

how to do something, neither is an agent. An agent's gonna have a really hard time with it. And then the third filter is, what is the worst-case scenario? So if your agent royally screws up, what happens, right? Sending a typo-filled social post, for example, that's not catastrophic. Sending a client the wrong contract? That's a problem. So using this framework we've been talking about here today, let's actually build something that will save you hours every week in your business. We all spend a lot of time, for example, in

our email, or even we might have

00:05:32,574 --> 00:06:02,644 [speaker_0]

somebody on our team doing our emails. So let's build an AI agent that handles your email support, and I'm gonna use an app called Relay.app for this demo, because it's beginner friendly, it's super intuitive. This is no code, doesn't require any coding. You could also build this in Zapier or Make or Lendi.ai or N8N. All right, so here we are inside of Relay.app. For this video, I wanna take you through actually building out the steps ourselves. When- wherever we're talking about an AI agent workflow, we want to consider, what do we want the trigger to be?

00:06:02,656 --> 00:06:32,654 [speaker_0]

So the trigger is just basically what wakes up our agent workflow. This can be manual, meaning you can come in here and press a button, and then it will trigger the workflow, or you can set it to trigger on a time, so every Friday at 8:00 AM. You can set it to trigger on an action, and that's what we wanna do. We wanna say add trigger. We're gonna click on Apps, and we're gonna click on Gmail, and we wanna do Email Received. So the trigger is whenever we get a new email into our Gmail account, it's gonna trigger the workflow. All right, so

00:06:32,694 --> 00:07:02,805 [speaker_0]

now that we've got the trigger, the next step is where we add the intelligence, and this is where we are connecting it to the brain, if you will, or we're choosing an AI model. So we're gonna click on Add Step, and in Relay here, that means adding AI here as to the next step. Now this is specific to Relay. They already have Classify as an option, so we're gonna click on Classify, and we wanna classify our email, so click on Next. One important thing here as you're building this, uh, this type of agent workflow out is you wanna figure out

00:07:02,815 --> 00:07:32,755 [speaker_0]

beforehand what type of email classification do you want to have. So for example, for me, I get lots of AI news and product updates. I want one for general types of questions that I get. I get a lot of collaboration emails. I get personal emails, like everybody else. Other, where we don't really know w- what category this email is gonna fit into. I've written a pretty basic instruction here. "You're expert email classifier for support system. Read the email carefully and classify it." The other thing

00:07:33,055 --> 00:08:03,076 [speaker_0]

is I get to choose what AI model I want to use. In this case here, uh, it defaults to GPT-5. And the next step is super simple, because I've clicked on Auto generate over here, and it's already taking the categories for the output. So if I mouse over Choices, you can see here, these are the categories that I've created over in my instructions over here. The next step is going to be, in Relay they call it flow control, so we're gonna create some paths. And remember that different

00:08:03,096 --> 00:08:33,096 [speaker_0]

email categorizations is going to be a path. So for example, the first path is gonna be

that first category of email that I designated, so AI news and product updates. So I can just create the path name right here, and then I gotta give it some sort of rule, and I'm just gonna click on Add Rule. And when it is classifying as the AI news and product updates, then follow this path. So for example, when the AI classifies as the AI news and product updates, I don't really need it to do anything when we

00:08:33,115 --> 00:09:03,235 [speaker_0]

get news or a product update, so I'm going to actually end the run at that moment when it gets classified as AI news and product updates. So now I still have four other email categories I wanna set up, so I'm gonna repeat this process. And the next one we're gonna do is the general question. So that's a path name, and again, I need to set the path rules, so I'm going to say Rules Match and Add Rule, Category, General Question. Okay. So we need to give it some resources to work with. So the first resource I wanna give it is I wanna

00:09:03,255 --> 00:09:33,296 [speaker_0]

give it some more intelligence so that it can figure out what to do when this type of general, quote-unquote, "email" comes in. So this is where I'm gonna use another AI step, and I'm going to do a custom prompt, and this is the prompt that I've already written. So basically, "You're a helpful and empathetic support assistant at an online business that teaches online entrepreneurs," blah, blah, blah. So I'm giving the AI instructions on how I want it to act, and the role I want it to take in this workflow. Now, the next integration we need to include in this AI step is we need to connect

00:09:33,316 --> 00:10:03,255 [speaker_0]

it to a knowledge file or knowledge files. So we need to give it some resources to work with to be able to do that. And what I recommend that you do in your business is to have an FAQ document so that you're tracking all the questions that you get in your business. So question, answer, maybe a Google Sheet. That's what I have. So I'm gonna click on the Knowledge button here, and I'm gonna look for FAQs. Here it is right here. So I'm going to attach my Google Sheet, that is the Email FAQs and Responses. So again, when I get an email in

00:10:03,755 --> 00:10:33,806 [speaker_0]

and the AI deems it, if you will, as a general question, it's gonna follow these instructions here as an empathetic, uh, email support assistant. It's gonna look at the email, and then it's gonna reference, all by itself, my Email FAQs and Responses to be able to answer that question, hopefully. Now, the next thing you wanna do is you wanna choose which AI model, uh, that you wanna attach to this step. I've tested the different models, and I think Claude Sign-up For is the best, so I'm gonna click on Use That Model. Now, as

00:10:33,836 --> 00:11:03,914 [speaker_0]

I mentioned just a second ago, what happens if a question comes into my email that is not included in the Email FAQs and Responses? What do we do? So down here, this is specific to Relay.app, you can click on Human Review, and it says, "Should a person review the AI output?" So I'm gonna say, "Only if AI cannot generate a valid output,"

meaning there's no answer in my knowledge file that I've given it, and I've asked it very specifically here in my instructions, "If no relevant answer is found in the sheet, do not generate a reply. Instead,

00:11:03,995 --> 00:11:33,434 [speaker_0]

leave the output empty so the workflow can notify the team," and that's exactly what I'm setting up over here. So, "Should a person review the AI output? Only if the AI cannot generate a valid output." And then, how do I want to be notified? So what I'm gonna do here is I'm going to, um, do a Slack channel, for example, so I'm gonna have that, and then I can give it the message, right? But I'm going to, uh, address somebody on my team. "The following question came in at the inbox and is not currently answered in our FAQ

00:11:34,035 --> 00:12:04,155 [speaker_0]

Google Sheet."And then I can just click on the at symbol, and this is where I can have it include any elements of the email that I want. So I want it to include the subject line, who it is from. I'm gonna say the email, and then I'm gonna click on Done. So, that is the hardest part, if you will, and that was super easy to do. That's pretty much the most complex step of this entire agent workflow. So, now what happens if we get an email that the AI can't answer and I wanna put it

00:12:04,216 --> 00:12:34,645 [speaker_0]

in my drafts folder in Gmail? Because I don't w- I don't trust it enough yet. So, I'm gonna click on Apps. I'm gonna go to Gmail, and I'm going to Create draft reply, and I'm gonna send it as me. Again, selecting an email, so I'm just selecting the email that came in from step number one. And then for the body of the email, I want it to be the response that Claude generated in the previous step. So I can just click on, uh, Reply email body, but I'm gonna click on DONE because I'm done

00:12:34,676 --> 00:13:04,995 [speaker_0]

in this step. So remember, we've got, in this example here, we've got five different email categorizations. You might have three, you might have two, you might have seven, whatever it might be. Again, I'm naming it, same thing, setting up a rule for it, so Classify, Networking, or Collaboration. So again, I'm telling it only come down this path when it's been categorized as a networking or a collaboration type of email, and I wanted to extract specific information out of that networking or collaboration email. So I'm gonna click on Extract

00:13:05,035 --> 00:13:35,005 [speaker_0]

to, Extract from, well, I want it to extract from the email. I'm gonna click on Next. I have layered in the instructions. So this is the prompt. I'm telling it what I want to extract out of the email, and so it's gonna do that all by itself. So I wanna kick that information into a Notion database. So, let's set that up. So I'm gonna click on Apps or I can just start typing there, so I'm gonna click on Notion. So I'm gonna add a page to, uh, database, and I'm gonna select the database and then choose YouTube Collab Requests. It just says, "Fields to populate," so I can just click on

00:13:35,076 --> 00:14:05,336 [speaker_0]

Add Field, and then I can just start, uh, selecting each one of these here. So, kind of collaboration that they want, I can click on Sender Company, so I just start mapping each of these fields. Super, super easy to do. And then I click on DONE, and just like that, we've updated my Notion database. The next thing I wanna do is I wanna let the team know, my team know, let's send a Slack message. Again, so we're- this is where we're adding the integrations into our workflow. So I want it to send a message, uh, send it to a channel, and I'm gonna say,

00:14:05,375 --> 00:14:35,596 [speaker_0]

"Rick to review" for this one here. And the message I wanna send is, "There's a new YouTube collaboration request you should check out." And I'm going to choose the at symbol because I want it to show the name of the tool, so I'm gonna select Tool Name. Here's the Tool Name, Page URL, you can obviously put in whatever you want. Uh, I can do Show Embedded Previews, I don't need to do that. Click on DONE, and that step is complete. So there's three of the five email categories that we have, uh, added to Integrations, and I just wanna complete this process out for the

00:14:35,635 --> 00:15:05,696 [speaker_0]

remaining two email categories. Okay, the next category is when it's a personal email, so I'm gonna go back and connect it to the Slack app. So I'm gonna send a message, and I want it to send directly to me, and I want the Slack DM to say, "Here's a personal email that just came in." So I can type the at symbol, and which will pull up this little menu here, and, uh, let's just say it is From, the Name. Then I want the subject line, and then the next one is the actual email, so I'm gonna choose the body

00:15:05,716 --> 00:15:35,765 [speaker_0]

of the email. And that's it. It's gonna send me a Slack message, cool. Let's set up the last email category. And the last category I'm gonna set up is Other. So, when an email comes in and it doesn't match one of these categories that we have given the AI to classify, so we'll give it the name Other, and then the rule for this is Follow this path when no other path is matching. And so, again, I'm gonna click on AI, and I'm gonna do Custom Prompt for this one here. And this is the prompt that I've writ-, "Your job is to carefully read the entire email and

00:15:35,775 --> 00:16:06,445 [speaker_0]

generate..." So it's generating, uh, basically a synopsis of what the email is. What do we wanna give it access to? So we wanna make sure that we give it access to the email that came in. So we've, it says, "AI has access to: EMAIL." We're using the GPT-5-mini model, and the output of this step is the Topic and Summary, as you can see right here. So again, the AI is figuring out what the topic is and the summary of whatever this email is about. For this purpose here, I'm just gonna set up a Human Review,

00:16:06,475 --> 00:16:36,466 [speaker_0]

where it says, "Should a person review the AI output?" Yes, Always. And then, again, I can s- I can have it emailed, Slack DM, Slack channel. I'm gonna say, "For Rick to review," and then, again, the message. So I can just write in here, "Please check the

following email. Here's a summary of what it's about." And that's it. I mean, it took a while to set it up just because I was talking through it. So the workflow triggers when an email is received in my Gmail. The intelligence, if you will, again, right here. So it's classifying the type of email that

00:16:36,576 --> 00:17:06,615 [speaker_0]

has come in, and then in the next step, this is where we've given it the in- integrations, we've given it the memory and, i.e. the knowledge files. We've added in the instructions on what to do when these types of emails come in, and it's all running autonomously while you focus on your higher value work. So the key mindset shift here when it comes to building AI agents is this. You're not building a tool. You're training a new team member. Give it context. Give it the necessary resources it needs to accomplish whatever task that you want it to do. The more

00:17:06,836 --> 00:17:36,996 [speaker_0]

specific your instructions, generally, the better your agent is gonna perform. Also, this is a very basic workflow. I can totally put truly agentic steps where I'm telling it, "This is what I want it to do," and I'm not giving it a whole lot of boundaries f- from which to work, and it just goes and figure out- figures out how to do it. So, start with one workflow, do your task audit, don't go too big too fast, and if you can use ChatGPT, you can start building AI agents in your business today. Thanks so much for watching the video today. Super appreciate it. I'll see you

00:17:37,036 --> 00:17:37,596 [speaker_0]

in the next one.