

What is Context Engineering?

Context engineering is the new buzzword,

right? But what if I told you it's actually an old idea, but just a new word for it. Now, you might think this is something as simple as giving your AI prompt the right context for whatever you're trying to do. And boom, you have 100% better output, but it's actually a lot more than that. This is a skill that you need to understand and get good at if you want to stay ahead with AI and truly create an AI first business. But the cool thing is you start implementing this in your workflows, you're going to be ahead of like 99% of people who are using AI. By the way, if you're new here, my name is Rick Mully and for the past 11 and a half years, I've run a multisefigure online business teaching other online business owners how to grow and scale their business. And today, I run an AI community for online businesses called the AI playbook where I help you leverage AI so that you can streamline your business, make it more efficient, and increase profit in the process. I'll link to it in the

Traditional Prompt Formatting vs. Context Engineering

description below. So, what is context engineering? Where traditional prompt formatting like we've always done in the past included a section for context, it tended to be static, right? Like this is my target audience or this is my writing style, right? In this style and it was done in chatbt or claude or Gemini, you know, one of those interfaces. But when it comes to context engineering, there are two primary differences. Number one, the context is dynamic in that you're connecting your prompt to either a dynamic knowledge base like a Google sheet or a Google doc that gets updated or you're connecting to a tool that you

use in your business like ConvertKit or HubSpot or what have you. And you're doing this through either an API or an MCP connection. And then number two, context engineering more pertains to workflows that contain AI within the workflows. for example, AI agents. So, let's say there are 10 steps in a workflow that you have and three of those steps are AI steps or AI agents.

Dynamic Context & Workflow Integration

Each of these steps would connect to a tool that you use in your business or a knowledge file for context and then use the output from the previous step if that's applicable. So, context engineering is actually referring to putting together the overall environment to accomplish something that you want to accomplish in your business, i.e. a workflow and the AI operates within this environment. You're trying to produce the highest quality outputs from your AI workflow with the right context with the right amount of context at the right time in the workflow. This is context

Context Engineering vs. Prompt Engineering with examples

engineering. Now, how is this different from prompt engineering? Well, prompt engineering is like a subset of context engineering. you're giving a set of instructions to an AI within the overall environment workflow, within the overall context engineering that you're putting together. Now, so many people essentially mocked the video that I did a year ago in July of 2024 talking all about stop learning prompt engineering. And this is exactly why because understanding what kind of context to give your AI is the key to getting the best outputs possible. Now, you still want good prompts, but they work within

the larger context framework or the workflow that you build. So, here's a perfect example of prompt engineering versus context and engineering. So, with a simple prompt, let's just say that you have an employee on your team and you say, "Hey, write a sales follow-up email to Sarah so that she becomes a client, right? It might be a decent email, but it lacks depth and any kind of contextual relevance and personalization to conversations that you may have had with Sarah. So, here's an example of what this overly simple, not very effective email might look like. Hey Sarah, hope this finds you well. Wanted to follow up on our recent conversation about our product services. Believe that our solution could be a great fit for your company's needs. Would you be available for a quick follow-up call this week to discuss further? Look forward to hearing from you soon. Super generic, lacks contextual relevance, lacks personalization, right? But that

Building Contextual Workflows (Relay.app Example)

is the prompt that we gave it. By the way, we could have put some context into that prompt following up to Sarah, but generally it would be static. Context engineering on the other hand is where you're building a workflow where the AI can access the prospect's company information. So Sarah's company's information, previous interactions with that person uh with your sales team, for example, specific pain points that are mentioned in previous calls, and maybe it's your CRM data before crafting a highly personalized follow-up email that the AI can do. So here's a perfect example of how this might look in an overall workflow. So here I am in relay.app and let's just say that you know it's triggered with a specific tag

for example um in my email CRM and that could be HubSpot or whatever it might be. And then the next step in my workflow is an AI step. So in this case here because I want to do writing I'm going to choose cloud signup 4 and I've given it a prompt. Okay. So yes I am giving it a specific prompt. However, the magic comes from number one, what I'm saying in the prompt because I'm telling it create a personalized outreach using the following context elements and I'm giving it what I want it to uh to include. And then down here, number one, I can include data from previous steps. So, data from the previous step in this workflow here. I can also connect it to tools. Now, this is where the context engineering, if you will, or the dynamic context comes in because I can connect this to any number of apps that are dynamically updated. I can connect this to an MCP server. I could build a custom tool in here. So, I'm getting personalized information that will help the AI create a contextually relevant, more personalized email response. Now I can also attach specific um knowledge files. So if it's a Google doc for example or a Google sheet and that means my knowledge file is dynamically updating. So this environment this workflow is staying contextually relevant because I'm connecting it to the right sources i.e.

Connecting Tools & Dynamic Knowledge Bases

context engineering. So for me this is how I look at context engineering. This is really about building AI workflows that contain AI agents because that's where things are going and agents in order for them to be effective. They need proper context. Now, could you do this in Claude, for example, where you can connect other tools? So, for

example, I can come in here and I did a whole video about Claude MCPs, which I'll link to right here, but you can do manage connectors. Also, they have this new feature here where they have these are all the pre-done MCPs that they have um that you can connect to really easily. Also, they have desktop extensions, too. So, if you're using Air Table, for example, and you are asking Claude to write an email to some person that's in your database in Air Table and you have a whole bunch of relevant context within Air Table, I look at that as a simple example of context engineering. But for me, the best way to be leveraging context engineering is building AI workflows in a tool that you're comfortable with. I love relay.app like I'm showing you here on the screen, but maybe you use MindPal or maybe you use N8N. Something else to consider when you're thinking about context engineering is how much is enough quote unquote context to provide within a workflow. And you can have too

How Much Context is Enough? Too Much vs. Too Little

Context

much context. So having too much context means that you're providing the AI with an excessive amount of knowledge or information that's either irrelevant, redundant, or it's just too much. It's overwhelming to the AI for the specific task that you're trying to accomplish. And if you have too much context, it can result in higher cost, higher higher API cost. The AI might struggle to find what's relevant and all the information that you're giving it. And also, this is going to increase the possibility of hallucinations. If you're asking an AI to say draft a short social media post about a new offer that you have or a new product or whatever and you feed it your

entire 50page business plan, every client testimonial that you have in your business, all your financial reports, that is too much context, right? Most of that information is completely irrelevant to the immediate task of writing a short social media post and you're making the AI work way harder than it needs to and potentially diluting the focus on the, you know, the offer, your product's key benefits. It's just too much context that's not necessary. And on the flip side of that, you can have too little context. This is where the AI is going to be lacking the necessary information to understand the nuances of your request, your business, or your specific situation. It's kind of like the the earlier prompt that we that I showed you where just write this sales follow-up email to try to get Sarah as a as a client. There wasn't enough context to have it do a really good job. It's going to give you a generic answer. It's more likely to hallucinate. There's a lack of personalization, etc., etc., etc. So again, if you ask an AI, write a social media post and that's it without providing any information about the product in your business or the offer, your target audience, your brand voice, your brand colors, your goal, the AI is going to produce an overly generic post that likely is not going to resonate with your audience. So what's the right amount of context? What's that sweet spot? Well, I like to think of this in

The REAL Framework for Context (Relevant, Efficient, Accessible, Logical)

terms of the acronym real. R E A L. So the R stands for relevant. This is where you're only including information that's directly relevant to that current task. If it doesn't help the AI achieve that

specific task or that specific goal, it's clutter. It's too much. The E stands for efficient. So, this is where you want to provide information concisely and without redundancy. This means you're summarizing long documents that aren't necessary. You're extracting key points and you're ensuring the AI isn't bogged down by unnecessary data. It also implies speed and cost effectiveness because again the AI is more efficient in what it's doing so it's going to cost you less. The A in real stands for accessible and this is where the context should be readily available to the AI when and where it's needed. So dynamic retrieval from knowledge bases or tools like I showed you in relay rather than being statically included in a prompt. Okay. Now the L in real stands for logical.

The information should be structured and organized in a way that the AI can easily understand and process. I got to tell you that so many business owners like I have in my community for example, one of the first things that we look at is are their knowledge files structured appropriately? Are they organized in a way that makes it easy for the AI to understand whatever the knowledge file is about. So this could mean using clear formats, consistent terminology, or well-defined knowledge bases. And if you're like, I have no idea how to do that. Well, besides joining us inside my community, you can ask the AI, ask ChatGPT, how best should I structure this piece of knowledge so that it's easy for you to consume and analyze and

Structuring Knowledge Files for AI

use when using it as a knowledge file?

And it'll tell you, will context

engineering solve everything that's wrong with AI, like hallucinations and

incorrect sources and just making stuff up, etc. Unfortunately, no. Right? Because of the nature of how LLMs work, like predicting the next word, for example, hallucinations can still happen even when using really smart context engineering in your workflows. And poorly worded prompt instructions also, which again are still important, can also lead to misinterpretations by the AI. However, it does go a long way in minimizing hallucinations because you're providing controlled and most and the most relevant information to the AI to use, especially when you're combining that with being specific in your prompting. This is grounding the AI in

Will Context Engineering Solve Everything?

real data. And then finally, as a side note, you can totally get more advanced within the individual concepts within context engineering. So, two things, for example, two ideas that you might want to think about. Number one, context evolution throughout your workflows. Think about this as the context is getting smarter and more specific as it moves through each step of your workflow. Again, you're going to be doing something like that in Relay or Mind Pal or NHN. Another advanced idea to be thinking about too is creating a learning system. So the workflow that you're building such that it's improving the context over time. So this is where your workflow or your system that you're creating remembers what worked from prior workflow runs and automatically improves its context decisions. So you're giving it feedback based on the output that it's giving you and you're building that into your workflows. And if you'd like to learn more about that

13:59: Advanced Ideas

stuff, we talk about that in my AI

community called the AI playbook. Again, I'll link to it in the description below. So, as you can see, context engineering, it's a super important skill to, if this is your first sort of exposure to it and what it's all about. Got to start learning it, start testing it out, start employing it into your overall workflows because again, it's only going to become more important the more we get into using agents and and so forth. This is going to put you way ahead of most people using AI when you can properly use context engineering within your overall environments and workflows where you're using AI. Thanks so much for watching today's video and I'll see you in the next