

Most AI agents don't fail because of weak models. They fail because of the data behind them. More than 90% of enterprise data is unstructured. Things like contracts, PDFs, Word documents, emails, transcripts, images, audio, video, and so much more. Unlike rows in a database, this content can't be easily searched, queried or fed directly into a model. That's why less than

1% of enterprise data makes its way into generative AI projects today. And here's the key:

public data is already baked into foundation models, so the real differentiator for AI is unlocking and harnessing enterprise data. Caroline, what makes unstructured data so difficult to

leverage? The challenge with unstructured data is that it's scattered across systems, inconsistent

in format, and often full of sensitive information. So, handing it straight to an AI agent risks

hallucinations, inaccurate answers or even leaks. To cope, data engineering teams have relied on

tedious manual work, sifting through disparate documents, stripping out sensitive details and

stitching together custom scripts. This does not make our engineer happy. The process can take

weeks. But the landscape is changing. That's why today we'll talk about two essential concepts: unstructured

data integration, which transforms raw content into AI-ready datasets in minutes, and unstructured data governance, which ensures those datasets can be discovered, catalog and trusted.

Together, they enable reusable, unstructured pipelines alongside structured ones, unlocking a

goldmine of data to power new use cases and address the technical challenges of integrating

unstructured data into AI workloads. This makes our engineers' lives a lot easier. Let's start with

integration. Adrian, can you describe what that looks like in practice? Of course.

Integration is

about transforming messy, raw, unstructured data into structured, machine-readable datasets. Think

of it as extending the familiar principles of structured data integration to a new modality.

Like ETL pipelines for structured sources, unstructured data integration creates repeatable

pipelines that ingest, process, and prepare high volumes of content. Only this time it's documents, emails,

chats, audio and more. The result? Users can automate in minutes what previously required weeks of custom scripting and maintenance. Here's how it works: We first

ingest data from sources like SharePoint, Box, Slack, Filestores and more, using prebuilt connectors. We then transform using prebuilt operators for text extraction, deduplication, language annotation, personally identifiable information removal, chunking content into usable segments and vectorizing those segments into embeddings. We finally then load embeddings into a vector database where they fuel retrieval augmented generation or RAG, AI agents, document classification, intelligence search and more, all without requiring deep machine learning expertise.

So, something like this? Yes, exactly. But what happens if a document changes? Updates don't require rerunning the entire pipeline. Only the delta is captured and pushed downstream, keeping pipelines current at scale without costly reprocessing. And for security, native access control lists support prev ... preserves document-level permissions so users and agents only see what they're authorized to, ensuring compliance and trust throughout the pipeline.

Unstructured data integration is a game changer, but it is only the first step. True unstructured data management goes beyond just integration. We also need to understand the data and trust it.

Caroline, how does that work? Integration focuses on data delivery and usability, but governance is what makes unstructured data truly discoverable, organized and trustworthy. Just as structured data has long benefited from data governance solutions, we now have end-to-end governance designed specifically to address the complexities of unstructured data. Let's walk through the steps. First, we connect to unstructured assets across the enterprise using prebuilt connectors. We then extract key entities like names, dates, topics, transforming raw files into structured analyzable data. Next enrichment pipelines classify content, assess quality and add contextual metadata. Documents are tagged with topics, people or sentiment to make them easier to organize and interpret. Results appear in simple validation tables with configurable rules and alerts that flag low-confidence metadata, helping ensure accuracy and trust. Assets then move through workflows into a central catalog, improving organization and discoverability. With technical and contextual metadata in place, users can now search and filter intelligently

across all assets. And finally, data lineage tracks how documents move from source to target, providing full visibility, compliance and auditability. With this governance layer, data teams deliver reliable, structured datasets that enable accurate AI model outputs and ensure compliance. Adrian, can these two technologies, unstructured data integration and governance, be used together? Unstructured data integration and governance close the reliability gap by giving AI agents high-quality, contextualized domain knowledge. With embeddings stored in a vector database, agents retrieve precise information instead of guessing, fueling more accurate RAG, copilots and domain-specific assistants. But the power doesn't stop with AI. The same foundation supports high-value use cases such as analytics and reporting. Teams can mine customer calls for sentiment trends, scan contracts to track compliance risks, or analyze field reports to uncover operational insights, all without manually sifting through thousands of files. Caroline, how do you see this shifting the enterprise AI story? It's a huge shift. Reliable AI agents require more than just smart models. They require smart data pipelines. Integration makes the data usable, and governance makes it trustworthy. But together, they unlock the 90% of enterprise data that's historically been out of reach. And this isn't just about AI agents. It's about giving enterprises new visibility into unstructured content. That's how teams can transition AI projects from prototypes to scalable production-grade systems.